

Weishu Zhang

626-566-4310 | weszhang@berkeley.edu | [LinkedIn](#) | [GitHub](#) | [Website](#)

Available December 2026. Backend-leaning full-stack engineer building production agent systems and enterprise workflow runtimes.

EDUCATION

University of California, Berkeley

Berkeley, CA

Bachelor of Science in Electrical Engineering and Computer Sciences

May 2027

– GPA: 4.0/4.0

– Relevant Coursework: Efficient Algorithms, Network Science, Artificial Intelligence

TECHNICAL SKILLS

Languages: Python, TypeScript/JavaScript, SQL, Java, Go, C/C++

Full Stack: React, Next.js, Node.js, FastAPI, PostgreSQL, Redis, SQLite, REST APIs

AI & Agents: LangGraph, MCP, LLM APIs, agent runtimes, structured generation, PyTorch, vLLM, QLoRA

Infrastructure: Docker, Kubernetes, gVisor, OCI, OpenTelemetry, pytest, distributed jobs, CI/CD

EXPERIENCE

Amazon

Incoming Software Development Engineer Intern

Sep 2026 – Nov 2026

ASML

San Jose, CA

Software Engineering in Test Intern

2026 – Present

- Built RRI for daily SQA regression tracking, cutting 80% of repetitive review work by automating Confluence review pages, Jira metadata, qTest sync, email, and compute-job follow-up.
- Designed a Test Plan Runtime Agent that drafts plans from cases covering 20+ coverage types, so QA only authors a few exceptional cases instead of handwriting the full plan.
- Compressed test-plan turnaround from 2 weeks of manual writing to as little as hours of human review, with the agent executing approved jobs asynchronously and writing results back.

Ludus

Remote

Co-Founder & Software Engineer, Cloud Agent Infrastructure

2026 – Present

- Built a Kubernetes-native sandbox control plane that packs 15+ concurrent gVisor-isolated agent instances onto a single AWS m6i.2xlarge (8 vCPU, 32 GB) via bin-packing, warm pools, and suspend/resume.
- Hardened untrusted-code execution with gVisor/Kata RuntimeClasses, namespace RBAC, default-deny NetworkPolicies, and CPU, memory, and ephemeral-storage quotas.
- Standardized reproducible agent environments with immutable OCI images, PVC-backed workspaces, TTL garbage collection, and OpenTelemetry traces so idle sandboxes yield host CPU.

Berkeley Artificial Intelligence Research (BAIR)

Berkeley, CA

Undergraduate Research Assistant

Jan 2026 – Present

- Led expert task acquisition for the Agent Human Last Exam benchmark, recruiting 30+ domain experts and contributing to the collection of 80+ complex, cross-domain agent tasks.
- Built an LLM-powered proposal agent that converts rough expert ideas into structured benchmark tasks, supporting iterative refinement, constraint checking, and higher-quality submissions.

University of California, San Francisco

San Francisco, CA

Software Research Assistant

Aug 2025 – Present

- Built a production-oriented long-context clinical information extraction pipeline, increasing inference throughput by 3x and reducing VRAM usage by approximately 80% using vLLM, paged attention, quantization, and batched execution.
- Created a 4,000+-example domain dataset, fine-tuned Qwen models with QLoRA, and implemented validation and self-consistency checks that maintained 100% schema-valid structured outputs.

OPEN SOURCE

PyTorch | [torch.compile](#), [torch.export](#), [Dynamo PRs](#)

- Contributed production fixes to pytorch/pytorch across compiler and export internals, resolving eager-versus-compiled correctness gaps, tracing failures, and export regressions with targeted tests and maintainer-reviewed patches.

SELECTED PROJECTS

GradeView | [React](#), [TypeScript](#), [Python](#), [PostgreSQL](#), [Redis](#) [Live](#) | [Demo](#) | [Poster](#)

- Refactored a production learning-analytics platform used in Berkeley courses into containerized React, Node.js, and Python services; standardized PostgreSQL as the source of truth and added Redis cache prewarming to improve reliability and read performance.