

Weishu Zhang

626-566-4310 | weszhang@berkeley.edu | [LinkedIn](#) | [GitHub](#) | [Website](#)

Backend-leaning full-stack engineer building production agent systems and enterprise workflow runtimes.

EDUCATION

University of California, Berkeley

Berkeley, CA

Bachelor of Science in Electrical Engineering and Computer Sciences

May 2027

– **GPA: 4.0/4.0.** Coursework: Efficient Algorithms, Network Science, Artificial Intelligence

TECHNICAL SKILLS

Languages: Python, TypeScript/JavaScript, SQL, Java, Go, C/C++

Full Stack: React, Next.js, Node.js, FastAPI, PostgreSQL, Redis, SQLite, REST APIs

AI & Agents: LangGraph, MCP, LLM APIs, agent runtimes, structured generation, PyTorch, vLLM, QLoRA

Infrastructure: Temporal, durable workflows, Docker, Kubernetes, gVisor, OCI, OpenTelemetry, pytest, distributed jobs, CI/CD

EXPERIENCE

Amazon AI Studio

Incoming Software Development Engineer Intern

Sep 2026 – Nov 2026


ASML

San Jose, CA

Software Engineering in Test Intern

2026 – Present

- Designed the **runtime agent infrastructure** for Tachyon SQA: daily regression tracking and a durable **test-plan lifecycle** (draft → QA approval → async job execution → result writeback).
- Built a **human-governed self-evolution** loop: reusable session memory is staged into shared memory, audited by humans, then persisted into the harness.
- Implemented **Temporal-style durable execution**: deterministic workflows, activities, QA-approval signals, task-queue claiming, idempotent writes, event-history replay, and sandbox recovery.



Open Ludus

Remote

Co-Founder & Software Engineer, Cloud Agent Infrastructure

2026 – Present

- Designed a **Kubernetes-native agent sandbox control plane** using CRDs and controllers to reconcile isolated, stateful workspaces from reusable runtime templates, with warm-pool provisioning, suspend/resume, and TTL-based garbage collection.
- Hardened untrusted-code execution with **gVisor/Kata RuntimeClasses**, namespace-scoped service accounts and RBAC, default-deny NetworkPolicies, restricted pod security contexts, and CPU, memory, and ephemeral-storage quotas.
- Standardized reproducible agent environments using immutable OCI images, validated configuration and short-lived secret projection, repository bootstrap, and PVC-backed workspaces; instrumented provisioning and execution with OpenTelemetry traces and lifecycle metrics.

Berkeley Artificial Intelligence Research (BAIR)

Berkeley, CA

Undergraduate Research Assistant

Jan 2026 – Present

- Led expert task acquisition for the **Agent Human Last Exam** benchmark, recruiting **30+ domain experts** and contributing to the collection of **80+** complex, cross-domain agent tasks.
- Built an LLM-powered proposal agent that converts rough expert ideas into structured benchmark tasks, supporting iterative refinement, constraint checking, and higher-quality submissions.

University of California, San Francisco

San Francisco, CA

Software Research Assistant

Aug 2025 – Present

- Built a production-oriented long-context clinical information extraction pipeline, increasing inference throughput by **3x** and reducing VRAM usage by approximately **80%** using vLLM, paged attention, quantization, and batched execution.
- Created a **4,000+**-example domain dataset, fine-tuned Qwen models with QLoRA, and implemented validation and self-consistency checks that maintained **100% schema-valid** structured outputs.

OPEN SOURCE & PROJECTS

PyTorch | [torch.compile](#), [torch.export](#), [Dynamo](#) PRs

- Production fixes to **pytorch/pytorch** compiler and export internals: eager-versus-compiled gaps, tracing failures, and export regressions, with maintainer-reviewed tests.

GradeView | [React](#), [TypeScript](#), [Python](#), [PostgreSQL](#), [Redis](#) Live | [Demo](#) | [Poster](#)

- Refactored Berkeley course analytics into containerized React/Node/Python services, with PostgreSQL as source of truth and Redis cache prewarming.